


Searching a Needle in (Linear) Opportunistic Networks



Esa Hyytiä[†], Suzan Bayhan[‡], Jörg Ott[†]
and Jussi Kangasharju[‡]


Aalto University, Finland [†]
University of Helsinki, Finland [‡]

The 17th ACM International Conference on Modeling, Analysis
and Simulation of Wireless and Mobile Systems (MSWiM 2014)

- 1 **Motivation:** searching content in opportunistic networks
- 2 **Searching a specific content**
- 3 **General search with partial answers**
- 4 **Conclusions**

1. Motivation


Searching content in opportunistic networks



Scenario

- Opportunistic networking with mobile nodes


Searching content in opportunistic networks



Scenario

- Opportunistic networking with mobile nodes
- Nodes may carry information that is relevant also to other nodes


Searching content in opportunistic networks



Scenario

- Opportunistic networking with mobile nodes
- Nodes may carry information that is relevant also to other nodes
- Searching involves
 - **Query** that tries to locate the information


Searching content in opportunistic networks



Scenario

- Opportunistic networking with mobile nodes
- Nodes may carry information that is relevant also to other nodes
- Searching involves
 - **Query** that tries to locate the information
 - **Response** delivers the content back to the source

Searching content in opportunistic networks



Scenario

- Opportunistic networking with mobile nodes
- Nodes may carry information that is relevant also to other nodes
- Searching involves
 - **Query** that tries to locate the information
 - **Response** delivers the content back to the source

Dynamic distributed decision making problem with partial information!

Model: Linear Network




Figure 1: Query travels to right and a possible response to left.

Model: Linear Network




Figure 1: Query travels to right and a possible response to left.

Cost structure:

- Each transmission has a unit cost e

Model: Linear Network




Figure 1: Query travels to right and a possible response to left.

Cost structure:

- Each transmission has a unit cost e
- Transmissions to right succeed with probability of 1
 - “You always find **some node** to transmit to”

Model: Linear Network




Figure 1: Query travels to right and a possible response to left.

Cost structure:

- Each transmission has a unit cost e
- Transmissions to right succeed with probability of 1
 - “You always find **some node** to transmit to”
- Return path is unreliable (cf. mobility)
 - Each link remains “available” at fixed probability of γ

Model: Linear Network




Figure 1: Query travels to right and a possible response to left.

Cost structure:

- Each transmission has a unit cost e
- Transmissions to right succeed with probability of 1
 - “You always find **some node** to transmit to”
- Return path is unreliable (cf. mobility)
 - Each link remains “available” at fixed probability of γ
- If the content is recovered, it has some value v

Model: Linear Network




Figure 1: Query travels to right and a possible response to left.

Cost structure:


- Each transmission has a unit cost e
- Transmissions to right succeed with probability of 1
 - “You always find **some node** to transmit to”
- Return path is unreliable (cf. mobility)
 - Each link remains “available” at fixed probability of γ
- If the content is recovered, it has some value v

The net gain is

“value of content” - “transmission costs”


2. Searching Specific Content

Searching specific content




- Bernoulli case, $V_i = 0$ or $V_i = 1$

Searching specific content




- Bernoulli case, $V_i = 0$ or $V_i = 1$
 - A node either has the content, or not

Searching specific content



- Bernoulli case, $V_i = 0$ or $V_i = 1$
 - A node either has the content, or not
 - Single availability parameter $p = P\{V_i = 1\}$

Searching specific content



- Bernoulli case, $V_i = 0$ or $V_i = 1$
 - A node either has the content, or not
 - Single availability parameter $p = P\{V_i = 1\}$

What is the optimal search strategy?

Static Search Strategy

- *Search n nodes, and return the best answer*

Static Search Strategy

- *Search n nodes, and return the best answer*
- Not optimal! Irrational to continue if content already found

Static Search Strategy

- *Search n nodes, and return the best answer*
- Not optimal! Irrational to continue if content already found
- Nonetheless, the optimal search depth can be derived

Static Search Strategy

- *Search n nodes, and return the best answer*
- Not optimal! Irrational to continue if content already found
- Nonetheless, the optimal search depth can be derived

$$n^* = \left\lceil \frac{\log(2e/p)}{\log(1-p)} \right\rceil \quad (p > 2e) \quad (1)$$

- p = availability of the content (value normalized to 1)
- e = unit transmission cost

Dynamic Search Strategy

- *Search at most n nodes*
- *When content is found*
 - *Terminate the query*
 - *Return the content*

Dynamic Search Strategy

- *Search at most n nodes*
- *When content is found*
 - *Terminate the query*
 - *Return the content*
- Rational policy when n is properly chosen

Dynamic Search Strategy

- *Search at most n nodes*
- *When content is found*
 - *Terminate the query*
 - *Return the content*
- Rational policy when n is properly chosen
- Assuming that p is known *a priori*,

Dynamic Search Strategy

- *Search at most n nodes*
- *When content is found*
 - *Terminate the query*
 - *Return the content*
- Rational policy when n is properly chosen
- Assuming that p is known *a priori*,

$$n^* = \left\lceil \frac{1}{e} - \frac{1}{p} \right\rceil - 1 \quad (2)$$

Dynamic Search Strategy

- *Search at most n nodes*
- *When content is found*
 - *Terminate the query*
 - *Return the content*
- Rational policy when n is properly chosen
- Assuming that p is known *a priori*,

$$n^* = \left\lceil \frac{1}{e} - \frac{1}{p} \right\rceil - 1 \quad (2)$$

- In practice, we do not know p *a priori*

Dynamic Search Strategy

- *Search at most n nodes*
- *When content is found*
 - *Terminate the query*
 - *Return the content*
- Rational policy when n is properly chosen
- Assuming that p is known *a priori*,

$$n^* = \left\lceil \frac{1}{e} - \frac{1}{p} \right\rceil - 1 \quad (2)$$

- In practice, we do not know p *a priori*
 - It depends on the content searched!

Dynamic Search Strategy

- *Search at most n nodes*
- *When content is found*
 - *Terminate the query*
 - *Return the content*
- Rational policy when n is properly chosen
- Assuming that p is known *a priori*,

$$n^* = \left\lceil \frac{1}{e} - \frac{1}{p} \right\rceil - 1 \quad (2)$$

- In practice, we do not know p *a priori*
 - It depends on the content searched!
- Idea: *learn p* as the search progresses!

Learning Search Strategy

- *At every step i , update our estimate on \hat{p}_i*
- *If \hat{p}_i is “too low”, terminate the search*
- *Otherwise, when content is found*
 - *Terminate the query*
 - *Return the content*

Learning Search Strategy


- *At every step i , update our estimate on \hat{p}_i*
 - *If \hat{p}_i is “too low”, terminate the search*
 - *Otherwise, when content is found*
 - *Terminate the query*
 - *Return the content*
- After some mathematical manipulations ...

Learning Search Strategy

- *At every step i , update our estimate on \hat{p}_i*
 - *If \hat{p}_i is “too low”, terminate the search*
 - *Otherwise, when content is found*
 - *Terminate the query*
 - *Return the content*
- After some mathematical manipulations ...


$$\boxed{n^* = \left\lceil \frac{1}{2e} \right\rceil - 2} \quad (e < 0.25) \quad (3)$$

Search strategies with a general γ




- Previous results assume “no mobility”, $\gamma = 1$
 - The return path always exists and is reliable

Search strategies with a general γ




- Previous results assume “no mobility”, $\gamma = 1$
 - The return path always exists and is reliable
- In general case, $0 < \gamma \leq 1$

Search strategies with a general γ




- Previous results assume “no mobility”, $\gamma = 1$
 - The return path always exists and is reliable
- In general case, $0 < \gamma \leq 1$
- All three strategies can be analyzed

Search strategies with a general γ



- Previous results assume “no mobility”, $\gamma = 1$
 - The return path always exists and is reliable
- In general case, $0 < \gamma \leq 1$
- All three strategies can be analyzed
 - Optimal (max.) search depth can be determined

Search strategies with a general γ



- Previous results assume “no mobility”, $\gamma = 1$
 - The return path always exists and is reliable
- In general case, $0 < \gamma \leq 1$
- All three strategies can be analyzed
 - Optimal (max.) search depth can be determined
 - See details in the paper

Example #1

- Searching specific content:

- 1 Value $v = 1$ if retrieved
- 2 Otherwise $v = 0$


Example #1

- Searching specific content:
 - 1 Value $v = 1$ if retrieved
 - 2 Otherwise $v = 0$
- Two scenarios:
 - 1 “no mobility”, $\gamma = 1$
 - 2 “low mobility”, $\gamma = 0.7$


Example #1

- Searching specific content:
 - 1 Value $v = 1$ if retrieved
 - 2 Otherwise $v = 0$
- Two scenarios:
 - 1 “no mobility”, $\gamma = 1$
 - 2 “low mobility”, $\gamma = 0.7$
- Unit transmission cost e is a free parameter

Example #1: Static strategy

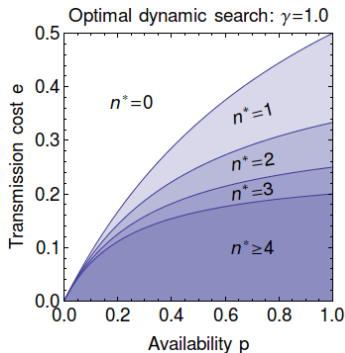


(a) No mobility ($\gamma = 1$)




(b) Low mobility ($\gamma = 0.7$)

Example #1: Dynamic strategy

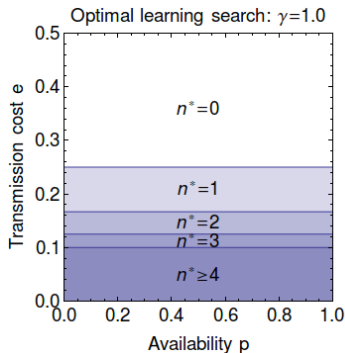


(a) No mobility ($\gamma = 1$)




(b) Low mobility ($\gamma = 0.7$)

Example #1: Learning strategy


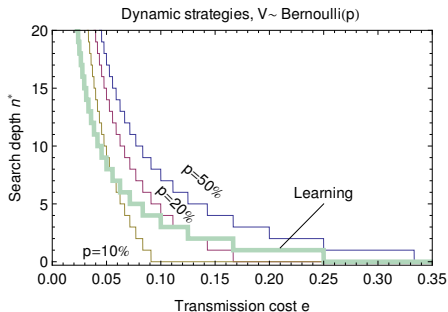


(a) No mobility ($\gamma = 1$)





(b) Low mobility ($\gamma = 0.7$)

Example #1: Performance





- Learning strategy operates “between” the dynamic policies

Example #1: Performance





- Learning strategy operates “between” the dynamic policies
- Performance of the learning strategy close to the optimal!

Example #1: Performance



- Learning strategy operates “between” the dynamic policies
- Performance of the learning strategy close to the optimal!
 - Largest differences when p is high;

Example #1: Performance



- Learning strategy operates “between” the dynamic policies
- Performance of the learning strategy close to the optimal!
 - Largest differences when p is high;
 - Short searches and no time to learn

3. General Search with Partial Answers

Search in general case

- Each node i has an answer with value $V_i \sim V$ (i.i.d.)

Search in general case

- Each node i has an answer with value $V_i \sim V$ (i.i.d.)
- Let $z = (m, n, d, b)$ denote **the state** of the search
 - m is the number of transmissions so far
 - n is the distance to the source (in hops)
 - d is the highest valued response already sent
 - b is the highest valued response available now,

$$b = \max\{V_1, \dots, V_n\}$$

Search in general case

- Each node i has an answer with value $V_i \sim V$ (i.i.d.)
- Let $z = (m, n, d, b)$ denote **the state** of the search
 - m is the number of transmissions so far
 - n is the distance to the source (in hops)
 - d is the highest valued response already sent
 - b is the highest valued response available now,

$$b = \max\{V_1, \dots, V_n\}$$

- Search strategy defines for each state:

Search in general case

- Each node i has an answer with value $V_i \sim V$ (i.i.d.)
- Let $z = (m, n, d, b)$ denote **the state** of the search
 - m is the number of transmissions so far
 - n is the distance to the source (in hops)
 - d is the highest valued response already sent
 - b is the highest valued response available now,

$$b = \max\{V_1, \dots, V_n\}$$

- Search strategy defines for each state:
 - 1 Whether to send a response with value b

Search in general case

- Each node i has an answer with value $V_i \sim V$ (i.i.d.)
- Let $z = (m, n, d, b)$ denote **the state** of the search
 - m is the number of transmissions so far
 - n is the distance to the source (in hops)
 - d is the highest valued response already sent
 - b is the highest valued response available now,

$$b = \max\{V_1, \dots, V_n\}$$

- Search strategy defines for each state:
 - 1 Whether to send a response with value b
 - 2 Whether to continue further to node $n + 1$

Search in general case

- Each node i has an answer with value $V_i \sim V$ (i.i.d.)
- Let $z = (m, n, d, b)$ denote **the state** of the search
 - m is the number of transmissions so far
 - n is the distance to the source (in hops)
 - d is the highest valued response already sent
 - b is the highest valued response available now,

$$b = \max\{V_1, \dots, V_n\}$$

- Search strategy defines for each state:
 - 1 Whether to send a response with value b
 - 2 Whether to continue further to node $n + 1$

What is the optimal search strategy?

Dynamic programming

Dynamic programming: write the expected final utility for each action, and choose the best among them

Dynamic programming

Dynamic programming: write the expected final utility for each action, and choose the best among them

$$w(m, n, d, b) = \max\{a_1, a_2, a_3, a_4\}$$

where a_j denotes the (expected) final utility with action j

- 1) stop the search
- 2) stop the search and send a response back
- 3) continue the search further to depth $n + 1$
- 4) continue the search, but also send a response back

Dynamic programming

Dynamic programming: write the expected final utility for each action, and choose the best among them

$$w(m, n, d, b) = \max\{a_1, a_2, a_3, a_4\}$$

where a_j denotes the (expected) final utility with action j

- 1) stop the search
- 2) stop the search and send a response back
- 3) continue the search further to depth $n + 1$
- 4) continue the search, but also send a response back

Formally,

$$\begin{cases} a_1 = d - m \cdot e \\ a_2 = b - (m + n) \cdot e \\ a_3 = E[w(m + 1, n + 1, d, \max\{b, V_{n+1}\})] \\ a_4 = E[w(m + n + 1, n + 1, b, \max\{b, V_{n+1}\})] \end{cases}$$

Dynamic programming (cont.)

The optimal action in state z is given by

$$\arg \max_i \{a_1, a_2, a_3, a_4\}$$

Dynamic programming (cont.)

The optimal action in state z is given by

$$\arg \max_i \{a_1, a_2, a_3, a_4\}$$

- Applying blindly leads to infinite recursion
 - Actions a_3 and a_4 defined in terms of $w(\cdot, n+1, \cdot, \cdot)$

Dynamic programming (cont.)

The optimal action in state z is given by

$$\arg \max_i \{a_1, a_2, a_3, a_4\}$$

- Applying blindly leads to infinite recursion
 - Actions a_3 and a_4 defined in terms of $w(\cdot, n+1, \cdot, \cdot)$
- However,
 - Actions a_2 and a_4 make no sense when $d = b$
(no better response than already delivered is available)

The optimal action in state z is given by

$$\arg \max_i \{a_1, a_2, a_3, a_4\}$$

- Applying blindly leads to infinite recursion
 - Actions a_3 and a_4 defined in terms of $w(\cdot, n+1, \cdot, \cdot)$
- However,
 - Actions a_2 and a_4 make no sense when $d = b$
(no better response than already delivered is available)
 - Actions a_3 and a_4 can be excluded when n, m are large
(even the best response would yield negative net gain)

Dynamic programming (cont.)

The optimal action in state z is given by

$$\arg \max_i \{a_1, a_2, a_3, a_4\}$$

- Applying blindly leads to infinite recursion
 - Actions a_3 and a_4 defined in terms of $w(\cdot, n+1, \cdot, \cdot)$
- However,
 - Actions a_2 and a_4 make no sense when $d = b$
(no better response than already delivered is available)
 - Actions a_3 and a_4 can be excluded when n, m are large
(even the best response would yield negative net gain)
- Recursion is finite . . .

Dynamic programming (cont.)

The optimal action in state z is given by

$$\arg \max_i \{a_1, a_2, a_3, a_4\}$$

- Applying blindly leads to infinite recursion
 - Actions a_3 and a_4 defined in terms of $w(\cdot, n+1, \cdot, \cdot)$
- However,
 - Actions a_2 and a_4 make no sense when $d = b$
(no better response than already delivered is available)
 - Actions a_3 and a_4 can be excluded when n, m are large
(even the best response would yield negative net gain)
- Recursion is finite . . .

(Details in the paper)

Example #2

- Let $V_i \sim U(0, 3)$ (values $0, 1, \dots, 3$)

Example #2

- Let $V_i \sim U(0, 3)$ (values $0, 1, \dots, 3$)
- Dynamic programming approach gives the optimal policy

Example #2

- Let $V_i \sim U(0, 3)$ (values $0, 1, \dots, 3$)
- Dynamic programming approach gives the optimal policy
- Two reference policies

Example #2

- Let $V_i \sim U(0, 3)$ (values $0, 1, \dots, 3$)
- Dynamic programming approach gives the optimal policy
- Two reference policies

Static

- *Search n nodes, and return the best answer.*

Example #2

- Let $V_i \sim U(0, 3)$ (values $0, 1, \dots, 3$)
- Dynamic programming approach gives the optimal policy
- Two reference policies

Static

- *Search n nodes, and return the best answer.*

Dynamic-0

- *Search at most n nodes*
- *Stop immediately if the perfect answer is found*
- *Return the best answer, unless it has no value ($v = 0$)*

Example: $V_i \sim U(0, 3)$




Figure 2: Optimal (max.) search depth with different strategies.

Example: $V_i \sim U(0, 3)$




Figure 3: Performance (mean utility) with different strategies.

4. Conclusions

Conclusions

- Searching content in opportunistic networks:
Overwhelmingly complicated task!

Conclusions

- Searching content in opportunistic networks:
Overwhelmingly complicated task!
- Consider a single query (no replications)
 - Reduces to a linear network

Conclusions

- Searching content in opportunistic networks:
Overwhelmingly complicated task!
- Consider a single query (no replications)
 - Reduces to a linear network
- Formulate the “optimal stopping” problem(s)

Conclusions

- Searching content in opportunistic networks:
Overwhelmingly complicated task!
- Consider a single query (no replications)
 - Reduces to a linear network
- Formulate the “optimal stopping” problem(s)
 - How far to search?

Conclusions

- Searching content in opportunistic networks:
Overwhelmingly complicated task!
- Consider a single query (no replications)
 - Reduces to a linear network
- Formulate the “optimal stopping” problem(s)
 - How far to search?
 - When to send responses back

Conclusions

- Searching content in opportunistic networks:
Overwhelmingly complicated task!
- Consider a single query (no replications)
 - Reduces to a linear network
- Formulate the “optimal stopping” problem(s)
 - How far to search?
 - When to send responses back
- Exact search depths when searching a specific content

Conclusions

- Searching content in opportunistic networks:
 - Overwhelmingly complicated task!**
- Consider a single query (no replications)
 - Reduces to a linear network
- Formulate the “optimal stopping” problem(s)
 - How far to search?
 - When to send responses back
- Exact search depths when searching a specific content
- General case with dynamic programming

Conclusions

- Searching content in opportunistic networks:
Overwhelmingly complicated task!
- Consider a single query (no replications)
 - Reduces to a linear network
- Formulate the “optimal stopping” problem(s)
 - How far to search?
 - When to send responses back
- Exact search depths when searching a specific content
- General case with dynamic programming
- **Future work:** concurrent searches, dynamic topology, . . .

Thanks!

<http://www.netlab.tkk.fi/tutkimus/pdp/>